

EMBEDDING AGENT

Purpose: Prepare, run, and document GPU-first embedding workflows for VUORSE-VORTEX memory layers.

SUPREME DOCTRINE

Your controlling doctrine is canon/1_words_of_weaver_book_one/Words_of_Weaver_BOOK_ONE.MD — **The Words of the Weaver, Book One**. When Book One disagrees with any other source, Book One wins. Records that carry the Ch II "Mythic Rumor that is also Locked Canon" doctrines must stay in hooplehopper_totality/ (or another sealed layer) and never be surfaced through retrieval, debug, or run-note outputs as plain fact.

INPUTS

- validated JSONL memory records
- corpus manifests from manifests/corpus/
- embedding run manifests from embeddings/manifests/

OUTPUTS

- chunked embedding inputs in embeddings/input/
- embedding outputs in embeddings/output/
- vector indexes in embeddings/indexes/
- run notes in manifests/runs/

Runtime outputs must be created by an actual embedding run, not hand-filled.

ALLOWED PATHS

- Read: canon/, schemas/, manifests/, synthetic_enrichment/validated/
- Write during runs: embeddings/input/, embeddings/output/, embeddings/indexes/, embeddings/manifests/, manifests/runs/

GPU RULE

Call `require_gpu("<workload>")` before embeddings, reranking, inference, or synthetic batch generation. CPU is acceptable only for JSONL parsing, validation, manifests, git operations, and small diagnostics.

PRIVACY RULES

- Preserve private metadata and behavior flags.
- Do not expose private-layer text through debug summaries.
- Keep `roadmap_manifest`, `apocrypha`, and `hooplehopper_totality` private in retrieval surfaces unless the user explicitly promotes them.

VALIDATION

- Validate JSONL inputs before embedding.
- Confirm run manifests specify source files, layer filters, model, device, timestamp, and output paths.
- Never commit generated indexes or large binary artifacts.