

DISTRIBUTED CONTINUITY



PURPOSE



Distributed Continuity prevents the Archive from depending on one record, one witness, one file, or one institutional memory. Slayverse continuity survives by being carried across documents, performances, rituals, indexes, artifacts, places, and embodied memory.

This folder describes decentralized preservation logic.

USE THIS FOLDER FOR



- Continuity maps across archive lanes.
- Notes on redundant preservation across canon files, rituals, performance backups, and indexes.
- Guidance for retrieval systems that need to find the same truth through multiple evidence paths.
- VUORSE-confidential continuity planning that should not become writer-room-facing canon.

CONTINUITY PRINCIPLES



- No single archive node should be treated as complete.
- Repetition across sources strengthens retrieval confidence, but does not override canon status.
- Contradictions must be flagged, not flattened.
- Private material can inform protection logic without entering public room knowledge.
- Distributed memory should clarify access, not blur permissions.

RECOMMENDED FIELDS

For machine-readable records, prefer:

- `id`
- `name` or `title`
- `type`
- `canon_status`
- `summary`
- `source_paths`
- `related_entities`
- `archive_lanes`
- `visibility`
- `confidence`
- `behavior`

REVIEW CHECKLIST

Before using a distributed continuity note downstream:

- 1 Confirm each source path exists.
- 2 Preserve canon labels from the source.
- 3 Separate locked canon from development or exploratory material.
- 4 Check whether the note reveals private finale or antagonist information.
- 5 For private synthetic outputs, enforce non-reveal and non-fact behavior flags.

The Archive survives because no single erasure can reach every place where memory has learned to stand.